

NCN/A3R Native Application Framework

("Native-Cloud/Native" services and
"Application-as-a-Resource")

NCN/A3R (hereafter **NA3**) is a **QT**-based application-development framework which prioritizes hybrid solutions combining cloud and desktop/native components. The **NCN** (Native-Cloud/Native) model refers to desktop client applications that are integrated with Cloud/Native back-ends; by sharing code libraries and data formats across both end-points, **NCN** solutions are more streamlined than native front-ends with generic back-ends, or Cloud/Native back-ends with web-application clients. The **A3R** (Application-as-a-Resource) model promotes self-contained, downloadable applications that can be distributed in source-code fashion and compiled with few (if any) non-**QT** dependencies. The combined **NA3** framework yields a comprehensive application-development toolkit with numerous components to streamline the implementation of **QT** applications (**NA3** can also be used as a template for implementations based on frameworks other than **QT**, such as wxWidgets or Operating-System-specific options).

The Current Status of QT Cloud Integration

There has been considerable demand in the native-application sector for a systematized Cloud Services model designed to interoperate with cross-platform native applications. Cloud/Native components can augment the functionality of native/desktop software by providing remote storage for user data; enabling users to share content for collaborative work; maintaining domain-specific repositories (i.e., spaces of resources whose format is specialized so that only select applications can access them properly); and upgrading or extending applications without re-install. We use the term "Native Cloud/Native" to describe hybrid applications whose server and client endpoints are both internally native – in contrast to conventional Cloud/Native where native servers are paired with (potentially) non-native clients.

Cloud/Native support in existing native-application frameworks is fairly primitive. Since **QT** is by far the most widely-used such framework, the **QT** case is instructive. In 2013 (following an earlier beta phase) the **QT** company introduced

Qt Cloud Services, which provided a convenient, Qt-aware cloud-hosting platform for **QT** accounts (in the company's words:

Qt Cloud Beta has solved an immense need for Qt developers when it comes to backend-as-a-service and believe that there is an even greater need to provide the Qt ecosystem with an all-in-one Qt solution for cloud computing). However — to the consternation of the Qt community — this project was discontinued several years later (retroactively we can identify some design flaws which might have hindered the project). Meanwhile, OpenShift discontinued their free-tier Cloud/Native hosting last year, and another company with Cloud/Native options, Arukas, is folding at the end of this month. This means that **QT** developers have limited options even for hosting hand-rolled **QT** cloud solutions (which can be done by compiling **QT** into a Ubuntu container)

Considering the prominence of both **QT** and Cloud/Native technologies in the contemporary computing landscape, it is disconcerting that no standard framework or hosting service provides a cloud platform which works with **QT** "out-of-the-box." The existence of such a platform would be a boon to software in sectors like scientific computing, bioinformatics, bioimaging, pharmaceuticals, academic publishing, and



other fields (where due to complex **GUI** and/or data-analytic requirements) the software is predominantly native-compiled and desktop-oriented.

Of course, many desktop applications have some web integration, but the current architecture forces the client-facing and web-facing components of the application to be almost completely separate, which adds to development time and expense. Moreover, current native-application environments do not fully leverage Cloud/Native services; they may well be implemented via more old-fashioned non-cloud servers. The great possibility of Native Cloud/Native is a peer-to-peer client-server relationship, sharing libraries and data formats on both ends; and the infrastructure to bring the benefits of Cloud Computing (e.g. faster development and deployment, and less expensive hosting, as compared to non-cloud web services) to the native-application sector.

Native Cloud/Native in the context of **NA3**

In light of the limitations just identified, LTS intends to contribute tools or hosting arrangements that would bring some of the capabilities of **QT** Cloud Services back to the market. The simplest commercial model for such a product is to licence a containerized **QT**-based **HTTP** server that can run as a local application during testing and development, before being deployed to a container hosting service. We have implemented a prototype server along these lines that we call NDP CLOUD (for "Native-Driven Platform"). NDP CLOUD is fully self-contained in a **QT** context (it bundles portions of the Node.js code base and utilizes the **QT** network module, so it requires no external **HTTP** or sockets libraries). One significant benefit of NDP CLOUD is that it is fully transparent: all of the code for parsing and routing **HTTP** requests can be loaded into **IDEs** (such as **QT** creator) and examined by the debugger. Another benefit is that project-specific libraries can be compiled into both NDP CLOUD instances and client front-ends; therefore, clients and servers can share procedures for serializing and deserializing domain-specific data structures. For development and prototyping, NDP CLOUD can be launched as an ordinary (non-virtual) **OS**; further testing can then be performed running NDP CLOUD as a local Docker container, before eventually deploying the application to a remote Docker hosting environment.

Via NDP CLOUD, **NCN** applications can be deployed on any Docker cloud service, such as OpenShift. In this guise LTS has no direct involvement with the hosting service (although NDP CLOUD includes some tools to streamline cloud deployment). Ideally, however, LTS would like to secure its own hosting capabilities, perhaps by using an LTS-specific container depoyed on OpenShift or a similar platform. LTS would allocate cloud assets to NDP CLOUD licensees (e.g. a limited free-tier hosting plan) for testing and development. A further possibility is to provide free hosting, subject to data-space constraints, to scientific institutions. The dwindling availability of free-tier Cloud/Native options is a hinderance to projects' adoption of Cloud/Native solutions for sharing and disseminating scientific data; this can result in researchers hosting data sets on platforms such as Mendeley or DataVerse, which have limited functionality or customizability compared to Cloud/Native containers. Use-cases for **NA3** in the context of scientific data sets are explained in the discussion of **A3R** below.

Application Development via **A3R**

The **A3R** model facilitates implementation of standalone native applications, whose data models and **UI** logistics are described via integrated metadata. As much as possible, **A3R** applications are entirely self-



contained, so that all application code and data can be packaged into single downloadable resource. In a **QT** environment, **QT** modules are available for concerns such as networking, database management, **C++** reflection, **XML** or **JSON** parsing, or embedded web viewers, so that **A3R** can leverage these capabilities without requiring separate library installs. For many use-cases, then, an entire desktop application can be deployed in source-code form, to be compiled and launched via a single click within **QT** creator. Self-contained in this manner, **A3R** applications can be treated as single resource units — to some degree analogous to Docker containers, but achieving their autonomy by leveraging the **QT** ecosystem rather than by virtualization.

A3R applications are also autonomous resources by virtue of detailed metadata bundled with application code. This metadata provides a summary of application-specific data models, capabilities, **UI** features, and user documentation. The metadata may be accessed by human users or by automated tools to help users become familiar with a newly-acquired **A3R** resource.

The **A3R** architecture is especially warranted when applications are designed to work with one or several non-standardized, domain-specific data formats (including those unique to an individual data set). In these scenarios, **A3R** applications provide both libraries for parsing and manipulating the domain-specific formats and a "reference implementation" documenting the proper visualization and User Experience optimal for the unique data structures involved. This structural profiling is advanced not only by domain-specific **GUI** components implemented within **A3R** applications, but also by **A3R** metadata which describes application-specific data types and interface requirements germane to any software components which work with such data types.

The **A3R** toolkit includes numerous components which may be useful to programmers implementing cross-platform, desktop-style applications, including a Hypergraph library for in-memory hypergraph-structured data, a built-in database engine for persistent, a parsing and grammar library, and a foundation for building customized scripting languages. These components have no external requirements and are distributed as raw **C++** files, that can be dropped in to any **QT/C++** project. As with NDP CLOUD on the server-side, these components are therefore "transparent": their code is directly bundled with application sources, and may be clearly examined in a debugging session. This is in contrast to typical libraries providing application-development features (such as database engines or source-code parsers), which typically require separate installation (often with separate build tools) and are opaque to the debugger. Another benefit of using internal **A3R** components is that they can be simultaneously compiled into **NCN** instances developed alongside. With that said, developers could certainly use non-**A3R** components in modular fashion (e.g., **QT**'s **SQL** data persistence) to replace their **A3R** equivalents (either initially or after an **A3R**-standalone prototyping phase).

NCN and A3R in consort

While **NCN** applications need not use **A3R**, or vice-versa, the two models are organically paired together. This can take the form of **NCN** servers hosting **A3R** applications as resources, and/or **A3R** software connecting to **NCN** instances as a domain-specific cloud back-end. The **A3R** metadata paradigm, based on "Hypergraph Ontologies", provides tools to streamline the encoding and distribution of application-specific data types. This model thereby accelerates the process of implementing cloud services procedurally aligned with **A3R** components, because complementary **A3R** and **NCN** endpoints can share the same



data-type libraries. Moreover, **A3R** interface definitions can serve as references for implementing compatible **NCN** server-side code; the interface specification documents which client-side procedures will handle any server-originating data structures, so the server-side data providers can be constructed accordingly.

To ensure rigorous alignment between client and server endpoints, **NA3** employs a data modeling paradigm based on hypergraphs; the mathematical framework for the relevant new hypergraph model (which adds some additional structure to the theory of Hypergraph Categories) is provisionally outlined in a chapter of a book which I edited that will soon be published (we can share this material, or alternatively a more thorough unpublished explication, as desired). In practical terms, the advantage of this model is that **QT**-specific data structures can be conveniently serialized and shared with or through cloud services; meanwhile the relevant data structures — and their interface and procedural requirements — are rigorously characterized, to support application testing, code-verification, documentation, and systematic User Interface development. **NA3** concretely operationalizes theories which have been advanced in the scientific-computing and knowledge-engineering community toward a more conceptually refined and “multi-scale” Semantic Web. **A3R** extends “Semantic Web alternatives” such as Conceptual Space Markup Language (**CSML**) and Categorical Informatics, while also providing a self-contained **C++** Hypergraph library comparable in some respects to AtomSpace (part of the OpenCOG platform) or to HypegraphDB (Hypergraphs and Conceptual Spaces have been proposed in combination as a comprehensive foundation for computational semantics, notably in the article *Interacting Conceptual Spaces* which arose from an Oxford University reading group on Category Theory and formal grammar).

A good example of an **A3R** use-case is that of hosting scientific data sets. According to the emerging “Research Object” paradigm, scientific data should be published alongside code which ensures that subsequent readers and researchers have the tools they need to access, analyze, and visualize the data set, including double-checking statistical analyses and/or replicating experiments. With **A3R**, data sets can be self-contained Research Object “bundles” while still being provisioned with full-featured desktop applications tailored to their specific information profile. LTS is actively developing a framework (called DataSet Creator, or **dsC**) for building data sets that paired with academic publications and with native software implemented in custom fashion for each data set (we can provide links to several data sets published as demonstration examples of this technology). The hosting and implementation of data sets along these lines offers concrete examples of **NA3** solutions and an opportunity to promote **NA3**-style development in the scientific and publishing communities.

As this use-case illustrates, our novel **NCN** model is based on rigorous serialization and interface-definition paradigms; in comparison, **QT** Cloud Services tended to reuse structures more appropriate for non-native contexts (e.g. **JSON**), which arguably limited client-to-server interoperability. LTS’s **NA3** model, by contrast, is combined with data modeling and serialization features that bring their own benefits to application projects over all; as such, this model does not only provide cloud-integration capabilities, but can be used as an overarching application-development framework.

